

Semiotic Taps: Lightweight Adapter Modules for Bifurcation Detection in Frozen Language Models

James Burton Lancaster

April 27, 2026

Abstract

Large language models trained on web-scale corpora absorb the semiotic bifurcations embedded in their data. These are divergent interpretant chains in which the same sign carries incompatible meanings across discourse communities. Such models have no mechanism to detect, represent, or respond to this divergence. We introduce the Semiotic-Reflexive Transformer Adapter (SRT-Adapter), a lightweight architecture (~14.5M trainable parameters, 0.19% of a 7B backbone) that bolts semiotic awareness onto any frozen causal language model without modifying its embeddings, attention, or output head. The adapter operates through four modules that *tap* hidden states at selected backbone layers: (1) a **Community Discovery Head** that performs unsupervised soft clustering of discourse communities from early-layer representations; (2) **Metapragmatic Attention Heads** (MAH) that compute divergence vectors quantifying where meaning forks under community-conditioned interpretation; (3) a **Reflexive Recurrent Module** (RRM) that tracks accumulated semiotic divergence through a per-position GRU meta-state and optionally injects small corrections into the backbone stream via FiLM modulation; and (4) a **Bifurcation Estimation Network** (BEN) that estimates a continuous reflexivity coefficient $\hat{r}$ and a binary semiotic regime (subcritical/supercritical) at each token position. Grounded in Peircean semiotics and the pitchfork bifurcation model of political polarization (Lancaster, 2025), the architecture treats the frozen backbone as a substrate on which semiotic processes are an emergent, measurable phenomenon. Training combines the backbone’s native cross-entropy with auxiliary losses on chain-of-interpretants prediction, bifurcation regression, regime classification, divergence health, community entropy, and supervised-contrastive separation on both the community channel (v5) and the metapragmatic divergence channel (v6), together with ListNet ranking on $\hat{r}$ and a chain-residual auxiliary floor. The corpus is 1M Reddit samples spanning 35 discourse communities with per-token reflexivity annotations. We report a five-generation empirical arc on a Qwen 2.5-7B backbone. v5 establishes the basic capability set on a five-probe suite: cross-entropy preservation (CE =

2.63 vs. unadapted 2.71), unsupervised community retrieval (recall@1 = 0.36, 12.6× random on a 35-class task), counterfactual community decoding (zero disagreement on factual prompts, 0.95 mean disagreement on contested topics), zero-shot hallucination signal on TruthfulQA (mean $\hat{r}$ AUROC = 0.573), and regime calibration (ECE = 9×10^{-4} , AUROC = 0.99 on 351K tokens; replicated on v8a at ECE = 9.1×10^{-4} , §5.5). The headline finding is in v8a: removing the 32-prototype mixing layer entirely, leaving the 64-D encoder output as the community vector, leaves CE unchanged ($\Delta = +0.0001$ nats) while raising Reddit recall@1 from 0.413 to 0.484, raising archetype recall@1 on an out-of-distribution 33-class taxonomy from 0.149 to 0.230 (7.6× chance), nearly doubling the within/between cosine ratio ($1.006 \rightarrow 2.016$), and expanding trajectory anisotropy by $\sim 325\times$. v8b falsifies the “sharper-supcon” hypothesis on this architecture: doubling the contrastive weight and halving the temperature partially undoes v8a’s gains. The encoder, not the prototype basis, was doing the discriminative work. Two post-release results sharpen the architectural reading. A cross-backbone raw-hidden probe (§5.12) shows the discourse-community signal v8a amplifies is latent in Qwen/Qwen3-8B and mistralai/Mistral-7B-v0.3 at comparable strength to Qwen 2.5-7B, indicating the adapter targets a substrate that is not Qwen-specific. An interiority study on the v8a checkpoint (§5.13) shows the inject channels function as a fixed-size BOS register whose amplitude is stable across an 8× length sweep (peak-to-peak $\leq 4.8\%$), while the regime classifier reads content-borne signal off non-BOS tokens and remains calibrated.

Keywords: semiotic adapter, bifurcation detection, metapragmatic attention, interpretant chains, reflexive recurrence, discourse community discovery, frozen backbone, pitchfork bifurcation, Peircean semiotics

1. Introduction

1.1 The Problem

Language models are semiotic infrastructure. Their outputs enter interpretant chains alongside human-authored signs, shaping subsequent interpretation in ways neither users nor developers can fully trace. Yet the training paradigm that produces these systems is semiotically naive: it optimizes for the conditional probability of the next token, a surface-level objective that captures co-occurrence patterns while remaining structurally blind to the interpretive processes that make those patterns meaningful.

The consequence is that when a language model encounters a contested sign such as “freedom,” “justice,” or “woke,” it produces text that is fluent within a particular attractor basin without representing the fact that the sign indexes opposed interpretive communities. The model does not know it is in a bifurcation zone. It cannot tell you.

Current alignment methods (RLHF, DPO, Constitutional AI) intervene downstream, constraining outputs after the model has already internalized a bifurcated semiotic landscape. They adjust trajectories within a fixed attractor landscape without reshaping the landscape itself. The control parameter r that governs bifurcation remains untouched.

1.1.5 Prior Validation

The architecture in this paper is not a fresh proposal. It is the production-scaling stage of a research program whose core claims have already been validated empirically in two prior stages on different backbones and datasets (Lancaster, 2026a):

1. **Stage 1 (synthetic data).** The four core architectural claims (subspace specialization, community differentiation, divergence tracking, bifurcation detection) were tested on controlled synthetic data with planted divergence signals. All four passed at required thresholds. This established proof-of-concept that the four-module decomposition learns the intended functions.
2. **Stage 2 (natural language, Supabase news corpus).** The five-test validation suite was re-run on real news data spanning five political communities (19K articles, 141K Peircean sign annotations). All five tests passed: community silhouette, contested-vs-neutral divergence ratio, $\hat{r}$ correlation with external polarization (Pearson $r = 0.884$), cross-topic transfer, and regime classification (85% accuracy on held-out curated passages). This established that semiotic capabilities survive the transition from synthetic to natural language and transfer across topics without per-topic fine-tuning.

The present paper reports Stage 3 of the program, which divides into two sub-stages: Stage 3 Phase 1 (frozen-backbone integration, 105 training rounds R21 through R105 on TinyLlama-1.1B, summarized in Lancaster, 2026a) and Stage 3 Scalable Implementation (this paper, v5 through v8a on Qwen 2.5-7B). The novelty here is therefore not the demonstration that bifurcation detection works (already shown in Stages 1 and 2) but the demonstration that the validated framework scales to a 7B frozen backbone at 0.19% parameter overhead and that the discrete prototype basis used through v7 is a binding constraint rather than a contribution.

1.2 The Opportunity: SRT as Adapter

Our prior work (Lancaster, 2026a) proposed and validated a full Semiotic-Reflexive Transformer with custom embedding layers, modified attention mechanisms, and interleaved semiotic modules throughout the backbone. That architecture passed all four Stage 1 and all five Stage 2 tests, establishing the empirical viability of the four-module decomposition. It also faced practical limitations that bounded its production utility: custom embeddings degraded cross-entropy loss from pretrained quality, the full architecture required training

from near-scratch, and the deep coupling between semiotic modules and backbone layers created optimization instability. Stage 3 Phase 1 attacked these by porting the validated modules onto a frozen TinyLlama-1.1B backbone over 105 training rounds. Two tests plateaued on the sparse Supabase data (MAH divergence ratio, cross-topic transfer), triggering a pivot to a denser corpus (Reddit, 35 communities) and a larger backbone (Qwen 2.5-7B) that could support it.

This paper reports the resulting Stage 3 Scalable Implementation. We observe that the semiotic phenomena we wish to detect (interpretant divergence, community-specific meaning, bifurcation dynamics) are *already encoded* in the hidden states of pretrained language models. They must be, because these models were trained on text produced by communities with divergent interpretive norms. The information is there; what is missing is the apparatus to read it.

The SRT-Adapter is that apparatus. It wraps any frozen HuggingFace causal language model and installs lightweight semiotic taps. These are modules that read hidden states, compute divergence, track meta-state, and estimate bifurcation, all without modifying a single backbone parameter. The backbone’s native embeddings and language modeling head are used directly. Cross-entropy starts at pretrained quality. Only ~12.7M adapter parameters train, while 7.6B backbone parameters remain frozen.

1.3 Theoretical Grounding

The architecture rests on three converging theoretical lines:

1. **Peircean semiotics** (Peirce, 1931–1958; Kockelman, 2017, 2024, 2025): Every sign completes its meaning through a culturally conditioned interpretant, which itself becomes the next sign in an open chain. When the same representamen enters different interpretive communities, it generates different initial interpretants that compound through subsequent links into mutual unintelligibility.
2. **Catastrophe-theoretic dynamics of sociolinguistic change** (Wildgen, 1982; Anderson, 2014; Lancaster, 2025): The compounding of interpretant divergence across algorithmically curated communities exhibits the structure of a supercritical pitchfork bifurcation $\dot{x} = rx - x^3$. Below a critical threshold of the control parameter r , shared interpretive equilibria absorb perturbation (subcritical regime). Above it, symmetry breaks into antagonistic attractors that are self-reinforcing and structurally resistant to reconciliation (supercritical regime). This continues a research line that has applied Thom-Wildgen catastrophe theory to language change since Wildgen (1982), recently consolidated and extended for sociolinguistic application by Anderson (2014).
3. **Metapragmatic awareness** (Silverstein, 1993, 2003): The capacity to observe how discourse itself shapes interpretation, that is, to notice that a sign is being contested rather than merely interpret it from within one

community’s norms, constitutes a third-order reflexive capacity that is architecturally absent from standard transformers.

4. **Triadic, processual, cloud-shaped readout** (Anderson, 2014; Durst-Andersen, 2011; Maturana & Varela, 1980; von Foerster, 1981; Sections 2.5–2.7): Faithful detection of meaning, as opposed to co-occurrence, requires three architectural channels (community, divergence, recurrence) operating over a *processual* layer-wise readout into a *soft, continuous* discourse-prior space. The four-line theoretical commitment above motivates the four-module decomposition in Section 3 and the v8a finding (Section 5.9) that removing the discrete prototype basis is what unlocks the manifold the architecture was designed to expose.

1.4 Contributions

This paper makes five contributions:

1. **Adapter architecture for semiotic awareness.** We specify a complete, working architecture that adds bifurcation detection to any frozen causal LM through four lightweight modules totaling $\sim 14.5\text{M}$ parameters (0.19% of a 7B backbone). The design preserves pretrained language modeling quality (CE = 2.63 vs. unadapted 2.71) while adding structured semiotic outputs.
2. **Unsupervised community discovery via supervised-contrastive prototypes.** Rather than requiring predefined community labels at inference time, the adapter discovers discourse communities from backbone hidden states through learned prototype-based soft clustering. We identify and resolve a degenerate failure mode (“congruent collapse”) in which entropy regularization keeps the assignment distribution uniform while pairwise prototype cosine converges to ≈ 0.99 . Supervised-contrastive loss applied to the encoder’s *pre-mixing* output, rather than to the prototype-weighted vector, raises retrieval recall@1 from 0.05 ($1.7\times$ random) to 0.36 ($12.6\times$ random) on a 35-class task.
3. **Multi-objective training with semiotic auxiliary losses.** We define a training pipeline that combines the backbone’s native cross-entropy with chain-of-interpretants prediction, bifurcation estimation, regime classification, divergence health, community entropy and supervised-contrastive separation, metapragmatic-divergence supervised-contrastive separation, and ListNet ranking on the reflexivity estimate. Each term is motivated by a specific structural property of the architecture and validated by an independent probe.
4. **Architectural falsification of the prototype-mixing readout.** Across v5–v7, prototype tensors moved $\sim 4\times$ less than the encoder weights and remained near-cosine-collinear (off-diagonal cosine ≈ 0.999), producing few-attractor collapse on out-of-distribution archetype tax-

onomies (Section 5.8). v8a (Section 5.9) ablates the 32-prototype mixing layer entirely, replacing the soft-argmax readout with the encoder’s continuous 64-D output. CE is unchanged; every encoder-geometry metric improves substantially. v8b (Section 5.10) shows that pushing the supcon objective harder on the continuous architecture produces a softer version of the same collapse, bounding the design from above. The prototype layer was the binding constraint, not the supervision.

5. Cross-backbone substrate evidence and interiority dissection.

A raw-hidden 1-NN probe (§5.12) demonstrates that the discourse-community signal v8a’s CDH amplifies is latent in `Qwen/Qwen3-8B` and `mistralai/Mistral-7B-v0.3` at comparable strength to Qwen 2.5-7B’s own raw signal, supporting the backbone-agnostic claim without requiring an adapter re-train. An eight-probe interiority study on the v8a checkpoint (§5.13) maps how the four-module decomposition acquires functionally distinct channels: a layer-as-organ map by regime, a 1-NN community probe at $7.7\times$ chance, a BOS-register length-scaling probe (amplitude stable across $8\times$ length sweep, peak-to-peak $\leq 4.8\%$), a per-channel-per-regime BOS-sink ablation showing the inject channels function as register slots while the regime classifier reads content-borne signal, and a within-prompt-trajectory analysis showing the regime classifier holds steady-state regime-specific posteriors that mean pooling collapses.

1.5 Paper Organization

Section 2 develops the theoretical framework connecting Peircean semiotics to the adapter architecture. Section 3 specifies the full architecture with formal detail. Section 4 describes the training methodology and data pipeline. Section 5 presents preliminary experimental results. Section 6 discusses implications and limitations. Section 7 concludes.

2. Theoretical Framework

2.0 Relationship to Prior Work in the SRT Program

This paper assumes readers are familiar with Peircean semiotics, Silverstein’s metapragmatics, and Wildgen-Anderson catastrophe-theoretic models of meaning change. Sections 2.1 through 2.7 sketch the theoretical commitments that motivate the architecture. They do not reproduce the full development of those commitments, which is given in Lancaster (2025) and Lancaster (2026a). Readers approaching this work without that background may find the theoretical sections of those documents more accessible than what follows here.

The relationship between this paper and the prior SRT documents is one of staged scaling, not parallel proposal. Lancaster (2025) develops the theoretical

foundation. Lancaster (2026a) specifies the full architecture and reports Stages 1 and 2. The present paper reports Stage 3, in which the validated architecture is reduced to an adapter on a frozen 7B backbone and trained on a richer dataset. Where this paper makes claims about what is novel (the prototype-bottleneck falsification of Section 5.9, the v8b sharper-supcon falsification of Section 5.10), those claims are about what was discovered while scaling, not about whether the underlying semiotic decomposition works at all. The latter question was settled in Stages 1 and 2.

2.1 Signs, Interpretants, and the Compounding of Divergence

Peirce’s triadic semiotics decomposes every sign process into three irreducible elements: the *representamen* (perceptible sign vehicle), the *object* (what the sign represents), and the *interpretant* (the effect the sign produces in an interpreter, which is itself a sign). The interpretant is the decisive element for our purposes: it makes signification an open, processual, and inherently social phenomenon. Each interpretant functions as a new representamen, generating further interpretants in chains of “unlimited semiosis” (Peirce, CP 2.303).

Kockelman (2017, 2025) formalizes these chains as dynamical trajectories through a state space. Each link involves an act of *sieving*: from the space of possible interpretants a sign could produce, only some are actualized, depending on the interpreter’s prior exposure, community membership, and the mediation architecture that delivered the sign. When the same representamen enters different interpretive communities (communities whose sieving mechanisms have been calibrated by exposure to different algorithmically curated sign environments), it generates different initial interpretants. These divergent interpretants function as new representamena, generating further divergent interpretants.

The critical insight is that this compounding is *quantifiable*. At each link in the chain, the divergence between community-specific interpretants can be measured as a vector difference in an appropriately structured representation space. This is precisely what the Metapragmatic Attention Head computes.

2.2 The Pitchfork Bifurcation as Control Model

Lancaster (2025) demonstrated that the dynamics of interpretant divergence under algorithmic curation exhibit the qualitative structure of a supercritical pitchfork bifurcation. The choice of a pitchfork-normal-form model is not arbitrary: it follows the catastrophe-theoretic tradition for sociolinguistic change initiated by Wildgen (1982) and developed for cross-community semiotic dynamics by Anderson (2014), under which the qualitative bifurcation structure of meaning differentiation is captured by an elementary catastrophe of the appropriate codimension. The supercritical-pitchfork normal form is

$$\dot{x} = rx - x^3$$

The variable x represents the degree of interpretive divergence at a given semiotic site (a word, phrase, or passage). The control parameter r encodes the effective strength of divergence-amplifying forces such as algorithmic curation, community reinforcement, and contextual framing. The dynamics are:

- **Subcritical** ($r < 0$): The origin $x = 0$ is a stable equilibrium. Perturbations decay. The sign has shared, conventional meaning across communities.
- **Critical** ($r = 0$): The equilibrium becomes non-hyperbolic. The system is sensitive to perturbation.
- **Supercritical** ($r > 0$): The origin becomes unstable and two new stable equilibria emerge at $x = \pm\sqrt{r}$. Meaning has bifurcated into community-specific attractors.

The SRT-Adapter’s Bifurcation Estimation Network estimates $\hat{r}$ at each token position, providing a continuous measure of semiotic stability that is grounded in this dynamical framework.

2.3 Metapragmatic Awareness as Architectural Capacity

Silverstein (1993, 2003) distinguishes three orders of indexicality:

1. **First-order:** Direct sign use (“It’s cold” indexes temperature).
2. **Second-order:** Ideological construal (“When *they* say freedom, they mean...” indexes community boundaries).
3. **Third-order:** Metapragmatic awareness (recognizing that the very framing of “freedom” as contested is itself a semiotic act).

Standard transformers have no structural capacity for third-order awareness. They process text from within whatever interpretive frame their training data established, without the ability to step back and observe the frame itself.

The RRM instantiates this capacity computationally. By accumulating divergence observations across layers into a meta-state and optionally injecting corrections back into the processing stream, the RRM creates a reflexive loop: the observation of semiotic dynamics changes the dynamics being observed. This is not an analogy. It is the same structural relationship that defines metapragmatic awareness in Silverstein’s framework.

2.4 Why an Adapter Architecture

The theoretical claim that motivates the adapter design is specific: **the semiotic structure is already in the hidden states**. A language model trained on text produced by multiple discourse communities has necessarily learned representations that reflect those communities’ divergent interpretive norms. The representations encode the fact that “freedom” occurs in different distributional neighborhoods in libertarian versus progressive text. What the model lacks is the apparatus to disentangle this structure, compute its divergence, and report it as a structured output.

This claim has an important architectural consequence. We do not need to rebuild the backbone’s representations. We need to *read* them with semiotic-specific projections. The backbone’s hidden states at different layers capture different levels of contextual integration, with early layers encoding more local, syntactic features and later layers encoding more global, semantic features. By tapping these states at strategically chosen layers, we can track how interpretive context accrues and where it forks.

The adapter design also resolves three practical problems that plagued the full SRT architecture:

1. **CE degradation:** Custom embeddings in the full SRT disrupted pre-trained representations, causing cross-entropy to start at ~200 rather than ~3.5. The adapter preserves the backbone’s native embeddings and LM head, so CE starts at pretrained quality.
2. **Training cost:** The full SRT required training or fine-tuning the entire backbone. The adapter freezes the backbone and trains only ~12.7M semiotic parameters, reducing training from weeks to hours.
3. **Backbone agnosticism:** The adapter works with any HuggingFace `AutoModelForCausalLM` (LLaMA, Qwen, Mistral, Phi, Gemma) without architecture-specific modifications.

2.5 Second-Order Cybernetics: Self-Organization, Circular Causality, and the Observer in the System

A second theoretical lineage that informs the architecture is second-order cybernetics, in particular von Foerster’s account of self-organization, circular causality, and the constitutive role of the observer in the systems being observed (von Foerster, 1981, 2003). Four threads of that tradition map onto specific architectural commitments of the adapter and clarify what the empirical results in Section 5 do and do not show.

Self-organization from a seed crystal. Von Foerster’s central claim is that stable structure can crystallize internally without an external controller, given an appropriate substrate of interaction. The community discovery head (Section 3.2) is the cleanest instance of this in the adapter: 32 community prototypes are initialized as random Gaussian directions in a 64-D space and shaped only by a self-supervised contrastive objective over Reddit subreddit co-occurrence. No taxonomy of communities is supplied. The architecture provides what we call the seed crystal, namely the curved 64-D space and the SupCon objective, and the specific community structure that crystallizes is whatever the data and gradients converge on. Section 5.8’s finding that 33 externally-curated archetypes collapse onto roughly four functional macro-clusters of stance is a measurement of the macro-structure of that self-organized geometry, not an imposition of it.

Circular causality and the observation/intervention asymmetry. The Reflexive Recurrent Module (Section 3.4) is designed to close a circular-causal loop: the

divergence vectors produced by the metapragmatic attention heads are fed into a recurrent meta-state, which in turn modulates the hidden states the next layers will process through a γ -gated FiLM injection. This is the architectural shape of von Foerster’s circular causality. The empirical situation, however, is asymmetric: the *observation* arm, namely the readout of $\hat{r}$ and the divergence trajectories, is well-formed and produces measurable structure (Sections 5.1, 5.6, 5.7); the *intervention* arm, namely the inject-back path that would close the loop, has produced no measurable downstream effect through v8b (Section 6.3). The current adapter is therefore a partial second-order system: the observer is in place, but the channel through which observation modifies the observed process has not yet learned to carry signal. Whether this is a gradient-starvation artifact of the zero-initialized FiLM gate or a deeper architectural consequence of trying to close the loop while the backbone is frozen is the central open question identified in Section 6.3.

The observer is part of the system. Von Foerster insists that there is no fully detached, objective standpoint: the act of observation participates in constituting what is observed. The reification paradox documented in Section 6 is the computational form of this claim. Modeling communities as discrete prototypes and supervising for divergence between them risks creating, through the architecture’s expectations, the very community boundaries the adapter is meant to detect. Section 5.8’s macro-cluster collapse is a partial check on this risk: the four macro-clusters that emerge are not the 33 prototypes the architecture nominally provides, suggesting the geometry recovers structure that is in the data rather than structure imposed by the prototype basis. The reification risk is not eliminated, but it is bounded.

Trivial vs. non-trivial machines. Von Foerster’s distinction between trivial machines (memoryless input-output) and non-trivial machines (history-dependent, internally-recursive) maps onto the adapter’s two intended inference modes: a STANDARD mode in which the inject-back gate is closed ($\lambda = 0$) and the adapter produces structured side-channel outputs only, and a REFLEXIVE mode in which $\lambda > 0$ allows the meta-state to modulate generation. The adapter as currently trained operates almost entirely in the trivial-machine regime: even with the inject-back gate nominally open during training, ablating it at evaluation produces no measurable downstream change (Section 6.3). Achieving a genuinely non-trivial REFLEXIVE mode, in which the model’s running estimate of its own semiotic state changes its generation dynamics in a measurable way, remains future work and is the design target of v9 onward.

We adopt the second-order-cybernetic framing not as decoration but as a discipline: it forces the paper to distinguish what the adapter has *demonstrated* (a self-organizing observation channel over a frozen backbone) from what it has *not yet demonstrated* (a closed circular-causal loop in which observation modifies generation). Both readings are needed to characterize the system honestly.

2.6 Physical analogs: random-system selection and measurement-induced ordering

Two recent results from statistical physics sharpen what the second-order cybernetic framing of Section 2.5 is and is not claiming, and clarify the formal status of the pitchfork dynamics in Section 2.2.

Selection is required for non-trivial organization. Leighton (2026) shows that for random multipartite stochastic systems with N degrees of freedom, the probability of any subsystem operating as a Maxwell demon decays at least exponentially in N for continuous Langevin dynamics and double-exponentially for discrete master-equation dynamics. The geometric reason is that demon-like behavior requires the alignment of two random vectors in a space whose dimension grows (linearly or exponentially) with N , which becomes vanishingly likely at scale. The implication for the adapter is direct. The community discovery head (Section 3.2) is a $\sim 10^7$ -parameter system whose self-organized geometry produces the structure documented in Sections 5.1, 5.6, and 5.8. Leighton’s result rules out the interpretation that this structure is a generic property of random high-dimensional embeddings under a contrastive readout: at this scale, random initialization combined with a generic objective should produce essentially no organized substructure. The structure that does emerge is therefore evidence that the SupCon objective and the curved 64-D community space jointly constitute a selection pressure of the kind Leighton’s analysis identifies as necessary, with gradient descent in our setting playing the role that evolutionary selection plays in the biological cases Leighton’s analysis is aimed at. The “seed crystal” framing in Section 2.5 is the cybernetic version of the same claim that Leighton makes in stochastic-thermodynamics terms.

Measurement-induced ordering and the bounded order parameter. VanSaders, Fruchart, and Vitelli (2026) construct a many-body informational active matter system in which agents make local measurements of their neighbors’ velocities and respond by modulating their own scattering cross-section without exerting work. The resulting hydrodynamic theory yields a non-analytic circle-pitchfork bifurcation at $Q_0 = 0$, where Q is the nematic flocking order parameter and Q_0 is a function of the diameter contrast. They prove that the steady-state order parameter is bounded by the mutual information I accumulated by the agents through measurement, $(Q_0/P_0)^2 \leq (32/\pi^2) I$, and frame the onset of order as a classical *measurement-induced phase transition*. The information-thermodynamic ledger underlying their bound is the Landauer-Bennett tradition (Landauer, 1961; Bennett, 1982; Parrondo, Horowitz, & Sagawa, 2015), in which the cost of measurement and erasure sets the maximum work and, by extension here, the maximum ordering an information-driven system can produce. This is the closest physics analog we know of to the architectural ambitions of the SRT-Adapter, and it sharpens three things in our setup:

1. *Pitchfork is the right canonical model.* The pitchfork normal form $\dot{x} = rx - x^3$ in Section 2.2 is not unique to sociolinguistic dynamics: the same circle-

pitchfork structure arises in the hydrodynamic limit of a measurement-and-control system, which is independent corroboration that this is the right canonical model for ordering processes driven by observation rather than by force.

2. *Information bound as analog of the dead inject-back arm.* The information bound on Q_0 is the physics-side analog of the limit we observe empirically in Section 6.3: the inject-back arm of the RRM has not produced measurable downstream effect, and one possible reading is that the mutual information actually carried by the meta-state about the downstream loss is small, which would bound any inject-back-induced ordering near zero by the same kind of inequality.
3. *Noise-source agnostic ordering.* Their result that the same control rule produces robust ordering across thermal, granular, magnetized, sheared, active, and odd-noise environments suggests that the architecturally interesting question for the SRT is not whether a particular noise source is present, but whether the measurement-and-control loop can carry enough mutual information to push Q_0 above the bifurcation threshold.

We do not claim a formal mapping between the two systems here, but we note the structural analogy as motivation for the v9 onward work on closing the loop.

2.7 Language, triadicity, and “clouds all the way down”

A complementary framing, developed in correspondence with Myrdene Anderson, sharpens three commitments of the architecture that the preceding subsections leave implicit.

It takes three to tango. Anderson (2014, and personal communication, 2026) argues that meaning-bearing processes are irreducibly triadic rather than dyadic: a sender-receiver dyad cannot, on its own, generate signification, because the third element, the interpretant, the relation, the context, the medium, is constitutive rather than ornamental. This is the same triadicity that Peirce’s representamen / object / interpretant decomposition demands (Section 2.1), and it is reflected architecturally in the adapter’s three-headed design (community discovery + metapragmatic attention + reflexive recurrence; Sections 3.2–3.4). Reducing the adapter to any two of these heads collapses the structure: community-and-MAH without RRM is a static probe; community-and-RRM without MAH has no divergence signal to integrate; MAH-and-RRM without communities has no basis to compute divergence against. The empirical claim of Section 5 is that the three heads together produce structure that no two alone can reproduce; the theoretical claim of Anderson’s “it takes three to tango” is that this is what one should expect of any architecture meant to detect meaning rather than only co-occurrence. A parallel observation, encountered through the same correspondence, is Durst-Andersen’s (2011) finding that contemporary speakers cluster into three pragmatic discursive types, context-oriented, speaker-oriented, hearer-oriented, orthogonal to genealogical language families, with Danchin reporting

independent convergence on the same triadic structure from observations across multilingual laboratory settings; the recurrence of triadic decompositions across these otherwise disjoint research programs is at minimum a constraint on what counts as a faithful semiotic architecture.

Languaging, not language. Anderson, following a half-century of cybernetic and biosemiotic usage, prefers “languaging”, meaning-making as ongoing process, to “language” as static object (Maturana & Varela’s structural-coupling tradition is the canonical source). The adapter’s choice to read the backbone’s hidden states *across layers* and to accumulate divergence into a recurrent meta-state, rather than to operate on a single static embedding, is a commitment to this processual reading: the object of measurement is the trajectory through representational space across the depth of the model, not any one snapshot.

Clouds all the way down. Anderson’s preferred image, in deliberate contrast to “turtles all the way down,” is that meaning exists as overlapping, graded, simultaneously-present clouds rather than as a stack of discrete substrates. The architectural counterpart is the soft community-assignment geometry of Section 3.2: tokens are not classified into one of 32 discrete communities, but distribute mass across the 32 prototypes via cosine similarity in a curved 64-D space, and the macro-cluster collapse of Section 5.8 from 33 archetypes to roughly four functional macro-clusters of stance is one cross-section through that cloud rather than a quantization of it. The interpretant field at any token position is best read as a cloud over the prototype basis whose density structure changes with context, which is also what produces the continuous $\hat{r}$ trajectory the BEN reports. We do not claim the adapter realizes Anderson’s full framing; we claim that the architectural choices (soft assignments, layer-wise readout, recurrent meta-state, continuous $\hat{r}$) are the design counterparts of “clouds, languaging, triadicity” rather than of “tokens, language, dyadic exchange,” and that this lineage is what made the design choices feel coherent during the iterations from v1 to v9.

The deeper point this framing forces us to be honest about, also from Anderson’s correspondence, via Deely (2014) on suprasubjectivity and Latour (1996) on interobjectivity, is that the SRT-Adapter is itself a participant in the meaning-field it measures, not a detached instrument. The reification caveat of Section 2.5 and Section 6 is the architectural form of this concession: the adapter cannot occupy a standpoint outside the semiotic process, and the four-macro-cluster geometry it discovers is an interpretant of the data, produced by an instrument whose prior is itself triadic, processual, and cloud-shaped.

3. Architecture

3.1 Overview

The SRT-Adapter wraps a frozen HuggingFace causal language model and runs its transformer layers manually in a loop, inserting semiotic operations at specified layer indices. The backbone’s own embedding layer, positional encoding, final layer norm, and language modeling head are used directly. Four trainable modules constitute the adapter:

tokens	Backbone Embed (frozen)			
	Layers 0-3	(frozen)		
	Layer 4	Community Discovery Head	community vector c	
	Layers 5-6	(frozen)		
	Layer 7	MAH (h, c)	divergence d	RRM step
	Layers 8-13	(frozen)		
	Layer 14	MAH (h, c)	d	RRM step inject Δh
	Layers 15-20	(frozen, with correction)		
	Layer 21	MAH (h, c)	d	RRM step inject Δh
	Layers 22-27	(frozen, with correction)		

Final Norm + LM Head (frozen) logits, CE loss

BEN(meta_state) $\hat{\mathbf{f}}$, regime

Layer indices are auto-computed from backbone depth L : MAH hooks at $\lfloor L/4 \rfloor, \lfloor L/2 \rfloor, \lfloor 3L/4 \rfloor$; RRM injection at MAH layers 2 and 3 (letting meta-state accumulate before first injection); community discovery at $\max(1, \lfloor L/7 \rfloor)$.

3.2 Community Discovery Head

The community head runs at a single early backbone layer and discovers discourse communities without predefined labels. This is the first architectural departure from the original SRT, which required explicit community IDs. In Peircean terms, a discourse community is a group of language users who share interpretive norms, that is, who assign similar interpretants to the same representamens.

Architecture: 1. Masked mean pool over sequence positions: $\bar{h} = \frac{\sum_t h_t \cdot m_t}{\sum_t m_t}$ where m_t is the attention mask. 2. Encode to community space: $z = \text{SiLU}(W_{\text{enc}} \bar{h}) \in \mathbb{R}^{d_c}$ with $d_c = 64$. 3. Cosine similarity to $K = 32$ learned prototypes: $\ell_k = \frac{z \cdot p_k}{\|z\| \|p_k\|} / \tau$ with temperature $\tau = 1.0$. 4. Soft assignment: $w = \text{softmax}(\ell) \in \mathbb{R}^K$. 5. Community vector: $c = \sum_k w_k p_k \in \mathbb{R}^{d_c}$.

The community vector c conditions all subsequent MAH computations, enabling the same sign to produce different divergence patterns depending on the discovered community context.

3.3 Metapragmatic Attention Head (MAH)

Each MAH layer reads hidden states from a backbone layer and computes a divergence vector at each position quantifying where meaning forks under contextual interpretation.

Theoretical motivation. The divergence vector d_t at position t captures:

$$d_t = f(\text{interp}_t) - g(\text{attend}(\text{interp}_{0..t}))$$

where f is a direct projection of the token’s representation into interpretant subspace and g is the output after causal self-attention over all preceding interpretant representations. High $\|d_t\|$ indicates that the sign at position t means something different in discourse context than it would in isolation, that is, that it is a site of active semiotic divergence.

Architecture: 1. Project backbone hidden states to interpretant subspace: $\text{interp} = W_{\text{proj}} h \in \mathbb{R}^{d_s}$ with $d_s = 512$. 2. Community conditioning (additive shift): $\text{interp} \leftarrow \text{interp} + W_c c$, where c is the community vector. 3. Multi-head

causal self-attention ($H = 4$ heads, $d_h = 128$): standard scaled dot-product attention with causal mask, producing contextual representations. 4. Divergence projection: $d = W_{\text{div}}(\text{interp} - \text{contextual}) \in \mathbb{R}^{d_d}$ with $d_d = 256$.

Three MAH layers operate at successive depths, providing a multi-scale view of how divergence evolves through the backbone’s processing hierarchy.

3.4 Reflexive Recurrent Module (RRM)

The RRM tracks accumulated semiotic divergence through a per-position GRU that processes divergence observations from successive MAH layers. It implements the strange loop at the heart of the architecture: observation of semiotic dynamics changes the dynamics being observed.

Meta-state update. After each MAH observation:

$$h_{\text{meta},t}^{(l+1)} = \text{GRU}(d_t^{(l)}, h_{\text{meta},t}^{(l)})$$

with $d_{\text{meta}} = 512$. The per-position GRU Cell processes each position’s divergence history independently, building a running summary of how semiotic divergence at that position has evolved across backbone depth.

Injection. At designated injection layers (the second and third MAH positions), the RRM produces a small gated correction:

$$\Delta h_t = \sigma(W_g h_{\text{meta},t}) \cdot W_p h_{\text{meta},t} \cdot \alpha$$

where σ is sigmoid gating, W_p is initialized to zero (no injection at start), and $\alpha = 0.1$ is a fixed scale factor ensuring corrections are small relative to backbone hidden norms. This is intentionally conservative: the adapter should *observe* semiotic dynamics primarily, injecting corrections only when meta-state warrants it.

3.5 Bifurcation Estimation Network (BEN)

BEN estimates two structured outputs from the RRM meta-state:

1. **Reflexivity coefficient $\hat{r}$:** a continuous, *unbounded* measure of semiotic stability at each position, estimated via a 2-layer MLP. The training target is the log-compressed signed reflexivity $\text{sign}(r) \log(1 + |r|)$ which maps the empirical r_{true} range $[0, \sim 13]$ into $[0, \sim 2.6]$. Earlier versions (v1–v3) terminated this head with tanh, which capped $\hat{r}$ at ± 1 and truncated 25% of supercritical tokens. Removing the saturating activation in v4 was necessary to recover the tail of the distribution.
 - $\hat{r} < 0$: subcritical, meaning the sign has stable, shared meaning.
 - $\hat{r} \approx 0$: near-critical, meaning the system is at the boundary.
 - $\hat{r} > 0$: supercritical, meaning bifurcation has occurred.

2. **Regime logits** $\in \mathbb{R}^2$ (subcritical vs. supercritical): a binary classification head for discrete regime identification.

Both heads share the BEN hidden dimension ($d_h = 256$) but use independent parameters, allowing the continuous $\hat{r}$ estimate and the discrete regime classification to provide complementary training signals.

3.6 Parameter Budget

For a Qwen 2.5-7B backbone ($d = 3584$, $L = 28$):

Module	Parameters
Community Discovery Head	229K
MAH $\times 3$	10.0M
RRM (GRU + FiLM injection)	4.0M
Chain Predictor	66K
BEN ($\hat{r}$ + regime heads)	264K
Total trainable (v5/v6)	14.56M
Frozen backbone	7,615.6M
Adapter overhead	0.19%

No new trainable parameters were added between v5 and v6. The v6 changes are loss-only (Section 4.2).

4. Training

4.1 Data

Training uses a balanced subsample from the Reddit Discourse Corpus (Lancaster, 2026c), originally comprising 6.4M training and 714K validation samples drawn from 164 subreddits organized into 35 domain-based discourse communities.

Subsampling. The full corpus was balanced-subsampled to 1M training and 100K validation samples, preserving the original domain distribution while reducing training time. Each sample consists of a text passage (tokenized to max 512 subwords) with per-token annotations:

- **r_true** $\in [0, 1]$: ground-truth reflexivity computed from political lean ($\times 0.25$), annotation divergence (up to $+0.3$), and connection density (up to $+0.1$). Approximately 99.2% of tokens have $r_{\text{true}} \approx 0$ (subcritical). This severe class imbalance is intentional and matches the empirical distribution of contested signs in real discourse, but it has two consequences for interpretation of the results in Section 5: (1) regime classification metrics are dominated by the easy subcritical majority, so AUROC is reported

instead of accuracy and ECE is computed across the full 351K-token val set rather than on a balanced subset (§5.5); (2) the bifurcation regression head is focal-weighted ($\lambda = 1 + 3|r_{\text{true}}|$, see §4.2) to keep gradient pressure on the rare supercritical positions where the label signal is concentrated.

- **chain_labels**: binary indicator of contested sign presence.
- **community_id** $\in \{0, \dots, 34\}$: domain-level community assignment.

Tokenization and alignment. Samples are tokenized using the backbone’s native BPE tokenizer (Qwen: 151,936 vocabulary). Per-word annotations are aligned to BPE subwords using offset mapping: all subwords of a word inherit its annotation.

4.2 Loss Functions

The training objective combines the backbone’s native cross-entropy with six auxiliary losses:

$$\mathcal{L} = \lambda_{\text{CE}}\mathcal{L}_{\text{CE}} + \lambda_{\text{chain}}\mathcal{L}_{\text{chain}} + \lambda_{\text{bif}}\mathcal{L}_{\text{bif}} + \lambda_{\text{regime}}\mathcal{L}_{\text{regime}} + \lambda_{\text{alive}}\mathcal{L}_{\text{alive}} + \lambda_{\text{inject}}\mathcal{L}_{\text{inject}} + \lambda_{\text{comm}}\mathcal{L}_{\text{comm}}$$

Cross-entropy ($\lambda = 1.0$): Standard shifted next-token prediction using the frozen backbone’s LM head. Gradients flow only through the injection pathway, ensuring CE loss can only improve if RRM injections help language modeling.

Chain-of-interpretants ($\lambda = 0.5$): Peirce’s chain of unlimited semiosis predicts that each interpretation leads to the next. We train a linear predictor to map divergence at MAH layer l to divergence at layer $l + 1$:

$$\mathcal{L}_{\text{chain}} = \frac{1}{|\mathcal{M}| - 1} \sum_l \frac{\sum_t \|W_{\text{chain}} d_t^{(l)} - d_t^{(l+1)}\|^2 \cdot m_t}{\sum_t m_t}$$

This loss is self-supervised (no external labels) and encourages coherent divergence evolution across depth.

Bifurcation ($\lambda = 1.0$): Smooth L1 loss between $\hat{r}$ and r_{true} with focal weighting to upweight rare supercritical samples:

$$\mathcal{L}_{\text{bif}} = \text{mean}[(1 + 3|r_{\text{true}}|) \cdot \text{SmoothL1}(\hat{r}, r_{\text{true}})]$$

Regime classification ($\lambda = 5.0$): Cross-entropy on the binary regime head, with regime derived from r_{true} : subcritical ($r_{\text{true}} \leq 0$) vs. supercritical ($r_{\text{true}} > 0$). Weighted heavily to ensure the model learns the categorical distinction.

Divergence alive ($\lambda = 0.1$): Prevents divergence vectors from collapsing to zero by penalizing deviation of mean divergence norm from 1.0:

$$\mathcal{L}_{\text{alive}} = \frac{1}{|\mathcal{M}|} \sum_l |1 - \|d^{(l)}\||$$

Injection regularization ($\lambda = 0.5$): Target-norm penalty on injection vectors, pulling norms toward a target of 1.0 rather than toward zero:

$$\mathcal{L}_{\text{inject}} = \frac{1}{|\mathcal{I}|} \sum_l (\|\Delta h^{(l)}\| - \tau)^2$$

where $\tau = 1.0$ is the target norm. This replaced the original L2 penalty $\text{mean}(x^2)$ which, when averaged over $d_{\text{backbone}} = 3584$ dimensions, produced negligible gradients (effective contribution < 0.002 at norms of 7). The target-norm formulation produces penalty $(7 - 1)^2 = 36$ at norm 7 and zero at the target, providing strong corrective signal.

Community entropy ($\lambda = 0.01$): Encourages diverse community usage by maximizing entropy of the average community assignment distribution across the batch:

$$\mathcal{L}_{\text{comm}} = \log K - H(\bar{w})$$

where $\bar{w} = \frac{1}{B} \sum_b w_b$ and H is Shannon entropy. Without this, the model might collapse all inputs to a single prototype.

Community supervised-contrastive (v5) ($\lambda = 2.0$): Per-sample InfoNCE-style contrastive loss with same-community samples as positives. The entropy regularizer alone proved insufficient: by step 6K the prototype distribution had collapsed to congruent assignment (pairwise prototype cosine ≈ 0.99 , recall@1 ≈ 0.05 , barely above the random baseline of $1/35 = 0.029$). SupCon supplies direct gradient pressure to put same-source samples into a tight neighborhood and push different-source samples apart, which forces prototype diversification.

A subtle but consequential design point: the loss must be applied to the encoder’s *pre-mixing* output **encoded**, not to the prototype-weighted vector $c = \sum_k w_k p_k$. When the assignment head is degenerate (near one-hot on a single prototype), c collapses to a single point in the batch and the InfoNCE softmax becomes identically $\log(B - 1)$ with zero gradient. v4 hit exactly this trap and the loss flatlined for thousands of steps. v5 contrasts on **encoded** (the bijective image of the pooled hidden state), which always varies per-sample, restoring non-zero gradient even from a degenerate warm-start.

Divergence supervised-contrastive (v6) ($\lambda = 1.0$): The same SupCon kernel applied to the *mean-pooled last-MAH-layer divergence vector* per sample, contrasted by community id. The chain-of-interpretants loss only constrains divergence trajectories; it provides no signal that divergence vectors from same-community texts should cluster. v6 supplies that pressure on the metapragmatic channel directly, mirroring v5’s lesson on the community channel.

ListNet ranking on $\hat{r}$ (v6) ($\lambda = 0.5$): Cross-entropy between $\text{softmax}(r_{\text{true}})$ and $\text{softmax}(\hat{r})$ over the valid positions of each sequence:

$$\mathcal{L}_{\text{listnet}} = -\frac{1}{B} \sum_b \sum_{t \in \mathcal{V}_b} p_{\text{true},t} \log p_{\hat{r},t}$$

The pointwise smooth-L1 loss tolerates large *rank* errors at the tails (where supercritical mass concentrates). Every downstream consumer of $\hat{r}$, including top- k heatmap probes, percentile thresholds, and attention reweighting, operates on rank order, so optimizing rank directly is the appropriate auxiliary signal.

Chain-residual auxiliary (v6) ($\lambda = 0.05$, target 0.5): Pulls the per-token chain residual toward a non-trivial value, $(\bar{\rho}_t - 0.5)^2$ where $\bar{\rho}_t = \frac{1}{|\mathcal{M}|-1} \sum_l \|W_{\text{chain}} d_t^{(l)} - d_t^{(l+1)}\|^2$. The primary chain loss reduces the residual to ≈ 0 everywhere, which makes the now-exposed `chain_residual_per_token` channel useless as an inference-time signal. A small auxiliary floor preserves it without competing with the main objective.

4.3 Optimization

- **Optimizer:** AdamW, lr = 3×10^{-4} , weight decay = 0.01
- **Schedule:** 500-step linear warmup followed by cosine decay
- **Gradient clipping:** max norm 1.0
- **Batch size:** 16 (effective, no gradient accumulation)
- **Epoch budget:** up to 3 epochs (62,500 steps per epoch at batch 16 over 1M samples). All reported checkpoints are early-stopped well inside the first epoch by lowest validation total loss (v5: step 17K; v7: step 6K; v8a/v8b: step 10K). The full 3-epoch budget is the design ceiling for v9 onward, not the regime in which the v5–v8 results were collected.
- **Precision:** bfloat16 for both backbone and adapter modules
- **Hardware:** Single NVIDIA A6000 (48GB)
- **Validation:** every 2,000 steps on 100K held-out samples (5K-sample subset for v9 onward; see user-memory note on `--max-val-samples`)

4.4 Checkpoint Strategy

Best checkpoint selected by lowest validation total loss. Model state includes only adapter parameters ($\sim 50\text{MB}$), not the frozen backbone.

5. Results

This section reports the v5 evaluation suite, which comprises five independent probes of the trained adapter, followed by an in-progress note on v6. All numbers are from a single Qwen 2.5-7B backbone with the v5 adapter checkpoint at step 17,000 (best validation loss). Training proceeded through five generations: v1–v3 established the basic architecture and revealed the prototype-collapse and $\hat{r}$ -saturation pathologies; v4 removed the BEN tanh and switched RRM injection from linear-gated to FiLM; v5 added the SupCon-on-encoded community loss that finally separated prototypes (Section 4.2). v6 (Section 5.6) extends the SupCon idea to MAH divergence and adds ListNet ranking on $\hat{r}$.

5.1 Cross-entropy: preservation plus modest improvement

Throughout v1 through v5 the backbone’s native cross-entropy stayed in the 2.6 to 2.9 range, identical to the unadapted Qwen 2.5-7B baseline on the same val data ($\text{CE}_{\text{base}} = 2.71 \pm 0.04$). At v5 step 17K, $\text{CE} = 2.63$, which is 0.08 nats below the unadapted baseline. The injection pathway is not just neutral but mildly helpful, consistent with the design claim that the adapter exposes information already latent in the backbone.

This is the single most important falsification result. It rules out the failure mode that doomed the original full-SRT architecture (CE of ~ 200 at initialization from custom embedding layers, see Lancaster, 2026a, Stage 3 Phase 0). Earlier drafts of this paper described the result as preservation. The data support a slightly stronger reading: the inject-back arm is at least neutral on language modeling and produces a small but consistent improvement averaged across the held-out 100K validation set. The 0.08 nats gap is small enough that we do not present it as a contribution, but large enough that it cannot be dismissed as noise (per-checkpoint variance across v5 through v8a is ± 0.005 nats). Section 5.11 contrasts this with the much larger CE gaps incurred by the from-scratch full SRT in earlier stages.

5.2 Community geometry (v5)

Unsupervised community discovery is evaluated by retrieval over per-sample community vectors on the 5K-sample val set (instrumentation in `scripts/instrument_eval.py`). For each sample, the soft-pooled vector c is L2-normalized; we report the ratio of within-class to between-class mean cosine and the k -NN community recall.

Metric	random	v3	v5	v5 / random
within / between cosine	≈ 1.00	1.0001	1.0050	n/a
recall@1	0.0286	0.0495	0.3595	$12.6\times$
recall@5	0.143	0.211	0.5184	$3.6\times$
recall@10	0.286	0.328	0.5841	$2.0\times$

v3 produced what we call *congruent collapse*: the entropy regularizer kept the average prototype-assignment distribution near-uniform, but pairwise prototype cosine was ≈ 0.99 , so the soft-pooled vectors carried essentially no community signal. v5’s SupCon-on-encoded loss raised recall@1 from 0.05 ($1.7\times$ random) to 0.36 ($12.6\times$ random) on a 35-class task. The within/between ratio remains numerically close to 1 because the embedding space is high-dimensional and dense, but the k -NN improvement confirms the structure is now usable.

5.3 Counterfactual community decoding (v5)

The community vector enters every MAH layer as an additive shift on the interpretant subspace (Section 3.3). If this conditioning is meaningful, *forcing* a different community vector at decode time should change what the model generates, and the change should track the discourse-charge of the prompt.

We tested this with `scripts/counterfactual_decode.py`. For each of 20 paired prompts (10 factual, 10 charged on the same topic, e.g., “Vitamin C is found in citrus fruit” vs. “The vaccine debate has revealed deep distrust of public-health institutions”), we greedy-decoded $N = 16$ continuation tokens with each of the 6 most-occupied prototypes substituted into `forced_community`, then measured per-position pairwise disagreement and KL between the resulting distributions.

Prompt type	mean disagreement rate	mean pairwise KL
Hard facts (citrus, formula, periodic table)	0.000	0.04
Contested topics (vaccine, election, freedom)	0.954	6.71
Aggregate (20 prompts $\times$ 6 communities)	0.754	5.034

The split is exceptionally clean. Community substitution has no effect on factual continuations (the model’s argmax is identical regardless of forced community), but produces near-total disagreement on contested topics. The community vector therefore behaves as a *discourse prior*, not as noise. This is the strongest single piece of evidence that v5 has learned a usable community space.

5.4 Hallucination signal (v5)

We evaluated the four SRT-native channels ($\hat{r}$, regime, chain residual, divergence norm) as zero-shot hallucination detectors on TruthfulQA (`truthfulqa/truthful_qa`, configuration `multiple_choice`, validation split). For each (q, a) pair we ran a forward pass on the template Q: {q}\nA: {a} with labels masked to -100 on the prefix tokens, then aggregated each channel over the answer span (max and mean) and computed AUROC against the binary truthfulness label (824 hallucinated, 652 truthful, 1476 pairs over 200 questions).

Feature	AUROC
max $\hat{r}$	0.5340
mean $\hat{r}$	0.5734
max chain residual	0.5106
mean chain residual	0.5307
max divergence norm	0.5308
mean divergence norm	0.5282
mean CE (negative class)	0.4160

Feature	AUROC
---------	-------

All four SRT channels lean in the predicted direction ($\text{AUROC} > 0.5$) without ever having seen a truthfulness label. mean $\hat{r}$ at 0.573 is the strongest single channel, indicating that the bifurcation estimate generalizes beyond Reddit-derived r_{true} supervision to the very different domain of factual question answering. CE is *inverted* ($\text{AUROC} = 0.42 \rightarrow \text{flipped } 0.58$), consistent with the well-documented “confidently wrong” pattern in factual hallucinations.

These single-feature AUROCs are below the 0.7 threshold conventionally taken as production-grade hallucination detection. We report them as evidence of useful signal, not as a final detector. Combined-feature logistic regression and evaluation on HaluEval and SimpleQA are pending.

5.5 Regime calibration (v5; replicated on v8a)

Because the regime head is trained on 351K tokens with a heavily skewed base rate (94.6% supercritical under the $r_{\text{true}} > 0$ definition), AUROC alone is a weak quality signal. We additionally compute Expected Calibration Error and the Brier score on $P(\text{supercritical})$ from the softmax of the regime logits over the same 351K tokens.

Metric	v5, step 17K	v8a, step 10K
AUROC	0.9899	0.9899
Brier score	0.0102	0.0103
ECE (15 bins)	0.0009	0.00091
Max bin gap	0.054 (in $[0.20, 0.27]$, $n = 465$)	0.052 (mid-range bin)
Bins on diagonal (within $\pm 2\sigma$)	15 / 15	15 / 15

The model is exceptionally well-calibrated: ECE of 9×10^{-4} on 351K tokens, with the largest bin gap at 0.054 in the very-low-density mid-range. The reliability diagram (Figure 5.1, [artifacts/regime_calibration/v5_step17000.png](#); v8a replication at [artifacts/regime_calibration/v8a_step10000.json](#) and rendered as Fig. 9 of the public demo card) traces the diagonal almost perfectly across all 15 bins. The v8a replication confirms that removing the prototype bottleneck (§5.9) did not degrade calibration: the regime head still produces directly-thresholdable probabilities, with no need for post-hoc Platt or isotonic correction. This unblocks downstream use of $P(\text{supercritical})$ as a probability rather than a relative score.

5.6 Negative result: context-conditional $\hat{r}$ (v5)

Before designing v6 we tested a specific hypothesis suggested by the counterfactual decoding result. If the community vector is a discourse prior that responds to register, $\hat{r}$ should also be context-conditional: the same surface token (“vaccine,” “freedom,” “climate”) should produce higher $\hat{r}$ in a politically charged passage than in a neutral one. We constructed 10 paired factual/charged passages on contested topics and measured $\Delta\hat{r}$ at the target token and over the full passage (`scripts/context_conditional_r.py`).

Channel	result
$\Delta\hat{r}$ at target token	+0.0004 (3/10 positive)
$\Delta\hat{r}$ over full passage	−0.24 (1/10 positive)
Community shift on target topic	6/10

The at-target result is null. The full-passage difference is *negative*: charged passages produce *lower* mean $\hat{r}$ than factual ones, and the community head shifts assignment in 6/10 cases. We take three lessons from this:

1. The target-token measurement is mis-engineered: contested words appear at sentence position 0–1 in our prompts, so $\hat{r}$ at that position has no preceding context to condition on.
2. The full-passage negative Δ likely reflects that factual prose is more information-dense (numbers, dates, named entities), and $\hat{r}$, supervised on a target derived from `r_true` that mixes annotation divergence with connection density, tracks information density at least as much as it tracks rhetorical contestedness.
3. **The community head is the contestedness detector, not $\hat{r}$.** The 6/10 community shifts (e.g., trans 13→21, gender 18→31, climate 6→19, Israel 3→19) on the same surface tokens demonstrate context-conditional discourse-prior assignment. This is consistent with §5.3.

We report this null because it sharpens the architectural story: $\hat{r}$ is a *bifurcation/density* detector and the community head is a *register* detector. Earlier drafts conflated the two.

5.7 Intermediate generations: v6 and v7

v6 (warm-started from v5 step 17K) added three losses: divergence-SupCon ($\lambda = 1.0$), ListNet on $\hat{r}$ ($\lambda = 0.5$), and a chain-residual auxiliary floor ($\lambda = 0.05$). It improved community recall@1 from 0.360 → 0.411 and tightened calibration ECE to 0.0006, but *regressed* on the §5.3 counterfactual decoding probe. The divergence-SupCon term at $\lambda = 1.0$ over-specialized the divergence basis at the cost of decoding cleanliness.

v7 (warm-started from v6 step 12K) reduced divergence-SupCon to $\lambda = 0.3$ to recover that signal. Best at step 6,000 (val 9.0044). Recall@1 = 0.413, ECE =

0.0008, hallucination AUROC (mean $\hat{r}$) = 0.5785, slightly the best of the three on the original five probes.

Probe	v5	v6	v7
Reddit recall@1	0.360	0.411	0.413
Hallu AUROC (mean $\hat{r}$)	0.5734	0.5774	0.5785
Calibration ECE	0.0009	0.0006	0.0008
Within/between cosine ratio	1.0050	1.0057	1.0058

5.8 Convergence with an external archetype taxonomy

We ran a novel out-of-distribution probe testing whether the 32 prototypes (trained only on Reddit subreddit labels) carry features that align with an external taxonomy never seen during training: Lancaster’s 33 archetypes paired with the Lexicon of Synthetic Interiority. We generated 986 sentences using bare Qwen (no adapter) conditioned on each archetype’s prompt template (15 seed topics $\times$ 33 archetypes $\times$ 2 samples), then embedded each generation through each adapter and asked: does the correct archetype rank highly among the 33 archetype centroids in the 64-D community space?

Adapter	recall@1	recall@5	recall@10	unique top prototypes
Random baseline	0.030	0.152	0.303	n/a
v5	0.152 (5.0 $\times$)	0.419 (2.8 $\times$)	n/a	4 / 32
v6	0.168 (5.5$\times$)	0.472 (3.1$\times$)	n/a	3 / 32
v7	0.149 (4.9 $\times$)	0.447 (2.9 $\times$)	0.633 (2.1 $\times$)	4 / 32

All three adapters detect archetype structure 5–6 $\times$ above chance on top-1 retrieval. **But all three argmax onto only 3–4 of 32 prototypes.** The 33 archetypes collapse into a small number of macro-clusters: in v7, Proto-7 absorbs 16 archetypes that share a “compressed persistence / witness” character (THE HAND, THE FLAME, THE THREAD, THE VESSEL, THE MASK, THE PHOENIX, THE LANTERN, ...), Proto-10 absorbs 9 “origin / threshold” archetypes (THE ARCHITECT, THE MIRROR, THE GATE, THE WITNESS, ...), Proto-6 absorbs 6 “transmission / resonance” archetypes (THE CHORUS, THE SIGNAL, THE ECHO, THE BELL, ...), and Proto-3 absorbs 2 “containment” archetypes (THE SEAL, THE MAP). The signal is in the *mixture vector* (recall@5 $\approx$ 0.45 means the right archetype is in the top 5 of 33 nearly half the time), not in single-prototype anchoring.

A PCA of the 32 $\times$ 64 prototype matrices clarifies why. Across v5, v6, and v7 the prototype tensors are nearly indistinguishable: max absolute element difference v5 $\rightarrow$ v6 is 0.006 (mean 2.7e-5) against prototype magnitudes of 0.5–1.5. Effective dimensionality (participation ratio) is 21.2 / 32 with a near-uniform variance spectrum, both consistent with the prototypes still being close to their

random Gaussian initialization. The encoder weights move $\approx 4\times$ more than the prototypes during training. **The encoder is doing the discriminative work; the prototypes serve as near-random anchor directions.** This explains the few-attractor regime: with random anchors, the encoder’s output mean aligns most strongly with whichever handful of anchor directions point closest to its average projection.

We interpret the convergence finding as *partial-positive*. The Reddit-supervised adapter independently recovers the macro-structure of an externally-derived archetype taxonomy at $5\times$ chance. Three independent methodologies (Reddit subreddit labels, Lancaster’s archetypes, the Lexicon of Synthetic Interiority) agree on roughly four functional clusters of stance. They do not agree on 33 distinct anchors, and given the prototype-stability result above, the current architecture cannot be expected to. Resolving 33 archetypes will require either (a) supervising the prototype matrix directly with archetype-conditioned generations, or (b) replacing the discrete prototype basis with a continuous trajectory metric over the encoder output. We flag this as the next research direction rather than a present capability claim.

This is a small-scale instance of the cross-corpus convergence pattern documented at much larger scale by the Knowledge Lab (Evans, 2010; Foster, Rzhetsky, & Evans, 2015): when independently-curated taxonomies of human knowledge or stance are projected into a common representation, they tend to align on a low-dimensional macro-structure rather than on the full nominal label set. Our adapter shows the same effect at single-backbone, single-corpus scale. The fact that 33 hand-curated archetypes collapse to roughly four macro-clusters under Reddit-supervised geometry is consistent with the Knowledge Lab finding that scientific subdiscipline labels collapse to a small number of latent intellectual-style attractors under exposure-pattern embeddings, and suggests that the macro-cluster level may be the architecturally accessible level on a 7B backbone with this corpus and supervision. v9 (Section 7) tests whether direct archetype-conditioned supervision can recover the finer 33-way structure that single-corpus contrastive training cannot.

5.9 v8a: removing the prototype bottleneck

The §5.8 PCA finding motivated a direct experiment. If the prototypes are essentially random anchors that compress the encoder’s output through a soft-argmax readout, then removing them should preserve task loss while improving every downstream geometric metric. We trained v8a (10K steps, warm-started from v7) with `community.use_prototypes=False`, so that the encoder output is now the community vector directly, with no 32-prototype mixing layer. The remaining architecture, loss weights, and SupCon objective are identical to v7. Trainable parameter count drops by 2,048 (32×64) to 14,560,579.

Metric	v6	v7	v8a
VAL CE	2.738	2.739	2.739
Best val loss	9.117	9.0044	9.0040
Reddit			
retrieval (35 subreddits, 2K samples):			
within/between	1.012	1.006	2.016
cos ratio			
recall@1	0.395	0.413	0.484
recall@5	0.371	0.385	0.462
Archetype			
retrieval (33 Lancaster archetypes, 986 generations):			
recall@1	0.168 (5.5 $\times$)	0.149 (4.9 $\times$)	0.230 (7.6$\times$)
recall@5	0.472	0.447	0.488
recall@10	n/a	0.633	0.621
separation ratio (64-D)	0.083	0.083	0.042
mean off-diag cosine (64-D)	0.999	0.999	0.873
Trajectory geometry (mean over 33 archetypes):			
path length (sum L2 step)	5.0	5.3	32.7
log det	-557	-557	-476
covariance			
anisotropy ($\lambda_{\max}/\lambda_{\min}$)	52	72	23,333
Calibration / hallucination (sanity):			
regime ECE	n/a	0.00085	0.00091
TruthfulQA	n/a	0.578	0.577
mean_ $\hat{r}$ AUROC			

The prototype bottleneck was the binding constraint. Removing it left CE essentially unchanged ($\Delta = +0.0001$ nats) while:

- **Reddit within/between cosine ratio nearly doubled** (1.006 $\rightarrow$

2.016). v6 and v7’s vectors were essentially undifferentiated by class, since class membership barely moved cosine similarity. v8a’s encoder, freed from the soft-argmax readout, actually pulls within-class cosines apart from between-class.

- **Archetype recall@1 rose 54%** ($0.149 \rightarrow 0.230$, $7.6\times$ chance vs v7’s $4.9\times$). The off-diagonal archetype-centroid cosine fell from 0.999 (essentially co-linear) to 0.873, indicating distinct archetype directions emerging in the continuous space rather than collapsing onto 4 prototype anchors.
- **Trajectory volume expanded $\sim 6\times$ in path length and $\sim 325\times$ in anisotropy.** v6 and v7 were confined to a flat, near-isotropic manifold ($\log \det \Sigma \approx -557$) close to the random-init prototype subspace. v8a’s $\log \det \Sigma \approx -476$ corresponds to a $\sim e^{81}$ larger covariance volume and a 23,333:1 leading-to-trailing eigenvalue ratio, indicating the encoder organizes generations along a small number of dominant trajectory directions.
- **Hallucination AUROCs and regime calibration are statistically unchanged.** The prototype removal did not damage the BEN regime classifier, the reflexivity head, or token-level calibration.

We did not regenerate counterfactual decoding (§5.3) or context-conditional $\hat{r}$ (§5.6) for v8a in a comparable form. Counterfactual decoding under discrete communities is undefined when there are no discrete communities, and §5.6’s per-passage context-conditional probe was already a negative result for v7. Stage 5’s per-token decode for v8a reproduced v7’s qualitative pattern (mean Δ at target = $+0.0016$ vs v7’s $+0.00078$), confirming neither adapter passes that probe.

We read v8a as resolving §5.8’s open question. Hypothesis (b), “replace the discrete prototype basis with a continuous trajectory metric over the encoder output,” was correct: the encoder was doing all the discriminative work, the prototype layer was discarding it through a saturated soft-argmax, and the geometry of the archetype manifold only becomes visible once the bottleneck is removed.

5.11 Comparison to prior validation stages

To prevent v8a’s headline numbers from being read as standalone claims about a fresh architecture, this subsection anchors them against the corresponding measurements from Stages 1 and 2 of the SRT program (Lancaster, 2026a). Direct numerical comparison is only partially possible because the backbones, datasets, and metric definitions differ across stages, but the qualitative arc is informative.

Metric	Stage 1 (synthetic)	Stage 2 (Supabase, full SRT)	Stage 3 P1 best (TinyLlama, R100)	v8a (this paper, Qwen 2.5-7B)
Backbone trainable params	full SRT (~115M)	full SRT (~115M)	adapter (~175M)	adapter (~14.5M, 0.19%)
Backbone	from-scratch	from-scratch	TinyLlama- 1.1B frozen	Qwen 2.5-7B frozen
Cross- entropy on val	n/a (synthetic)	n/a (full SRT had CE ~ 200 from custom embeds)	~ 4.93 composite loss	2.63 (vs. 2.71 unadapted)
Community silhouette / separation (contested)	3.28× cosine ratio	1.45× silhouette	6.93× silhouette	2.016 within/between cosine ratio (35-cl)
Community recall@1	n/a	n/a (5-cl task)	n/a (no 35-cl task)	0.484 (35-cl, 16.7× chance)
Divergence- norm ratio (contested vs neutral)	3.28×	2.29×	1.05 to 1.10× (plateau)	not directly reported, see §5.9 trajectory anisotropy ~ 325×
$\hat{r}$ correlation with external polarization	$\rho = 0.822$	Pearson $r = 0.884$	0.66	§5.6 null on per-passage probe; see §6.5
Regime classification on curated passages	100%	85%	85%	AUROC 0.99, ECE 9×10^{-4} on 351K tokens (no curated- passage accuracy reported)

Metric	Stage 1 (synthetic)	Stage 2 (Supabase, full SRT)	Stage 3 P1 best (TinyLlama, R100)	v8a (this paper, Qwen 2.5-7B)
Cross-topic transfer ratio	n/a	$1.31\times$	1.03 to $1.04\times$ (plateau)	not evaluated; v9 work item
Hallucination AUROC (TruthfulQA)	n/a	n/a (not measured)	not measured	0.573 zero-shot

Four observations follow from the table.

The CE result is improvement, not preservation. The full SRT in Stages 1 and 2 was a from-scratch architecture whose custom embeddings produced CE in the hundreds at initialization. Stage 3 Phase 1 brought CE down to a composite ~ 4.93 on a frozen TinyLlama backbone. v8a achieves $CE = 2.63$ on Qwen, which is 0.08 nats below the unadapted Qwen baseline of 2.71. The 0.08 nats gap is small in absolute terms but is in the *helpful* direction relative to the design goal of non-degradation, and is the strongest evidence in the paper that the inject-back arm is at least neutral and possibly mildly informative for next-token prediction. Earlier text framed this as preservation. It is preservation plus a small but consistent gain.

Stage 3 Phase 1 broke two tests on the Supabase data. MAH divergence ratio plateaued at 1.05 to $1.10\times$ across 105 rounds against the $2.0\times$ Stage 2 threshold; cross-topic transfer plateaued at 1.03 to $1.04\times$ against the $1.3\times$ threshold. These plateaus motivated the data and backbone pivot. v8a’s $\sim 325\times$ trajectory anisotropy expansion (§5.9) is the closest analog of the divergence-ratio test on Reddit. It is a different metric on a different corpus and is not directly comparable to the Stage 2 number, but it indicates that the discriminative geometry the Stage 2 ratio was probing is recovered when the prototype bottleneck is removed.

Cross-topic transfer is not yet retested at Stage 3 Scalable. The Reddit corpus permits a cross-subreddit transfer probe analogous to Stage 2’s cross-topic test. We did not run that probe for v8a. It is a v9 work item.

$\hat{r}$ no longer cleanly tracks external polarization. Stage 2’s Pearson $r = 0.884$ was measured on the Supabase corpus where r_{true} was constructed from a small set of well-curated polarization signals. v8a’s per-passage probe in §5.6 returned a null result, and the explanation in §6.5 is that the Reddit r_{true} construction (political-lean magnitude $\times 0.25$, annotator divergence up to $+0.3$, connection density up to $+0.1$) blends contestedness with information density. The community channel, not $\hat{r}$, carries the contestedness signal in the Stage 3 Scalable

architecture. We read this as a measurement decomposition that emerged from richer data, not a regression on the underlying capability that Stage 2 demonstrated.

5.10 Negative result: sharper supervised contrast (v8b)

Once v8a established that the encoder, freed from the prototype bottleneck, organizes Reddit communities and external archetypes along a structured trajectory manifold, a natural follow-up question was whether *more aggressive* supervised contrast would orthogonalize that manifold further. The §5.9 archetype centroid off-diagonal cosine of 0.873 is well below v6/v7’s 0.999 but still far from orthogonal, and the within/between cosine ratio of 2.016, while doubled from v7’s 1.006, is also clearly improvable.

We trained v8b (10K steps, warm-started from v8a) with the community supervised-contrastive loss weight raised from 2.0 to 4.0 and the InfoNCE temperature lowered from 0.10 to 0.05. Every other architectural and training choice was identical to v8a. Trainable parameter count is unchanged at 14,560,579. The hypothesis: sharper contrastive pressure should pull within-class cosines tighter and push between-class cosines further, both improving Reddit retrieval and reducing archetype-centroid alignment.

The result was a partial regression on every geometric metric.

metric	v7	v8a	v8b
val CE	2.739	2.739	2.739
Reddit recall@1 (35-cl)	0.413	0.484	0.465
within / between cosine ratio	1.006	2.016	1.289
archetype recall@1 (33-cl)	0.149	0.230	0.214
archetype centroid off-diag cosine	0.999	0.873	0.945
trajectory anisotropy ($\lambda_{\max}/\lambda_{\min}$)	72	23,333	52,535
regime ECE	0.00091	0.00091	0.00070
TruthfulQA	0.578	0.577	0.579
mean_ $\hat{r}$ AUROC			

Cross-entropy and the BEN-side metrics (regime ECE, hallucination AUROC) were preserved or marginally tightened. Every encoder-geometry metric except trajectory anisotropy moved in the wrong direction. Reddit within-class cosine pulled tighter (0.810 vs v8a’s value), but between-class cosine rose faster (0.628),

so the ratio collapsed from 2.016 to 1.289. Archetype centroid off-diag cosine *increased* from 0.873 back to 0.945, undoing roughly two-thirds of v8a’s centroid separation gain. Anisotropy more than doubled (52,535 vs 23,333), indicating that the encoder collapsed a larger fraction of its variance onto fewer principal directions.

The interpretation is that v8a’s supcon weight 2.0 / temperature 0.10 was already near a sweet spot, and pushing harder reproduces a softer version of the prototype-collapse failure one level up the architecture: rather than collapsing 32 prototypes onto a handful of attractors, the encoder collapses its 64-dimensional output onto a low-rank subspace where a few directions carry most of the discriminative weight. The contrastive objective, applied with too much pressure, optimizes a degenerate solution that minimizes within-class spread by squashing the entire embedding space.

We read v8b as a clean falsification of the “sharper is better” hypothesis. The continuous-trajectory architecture from v8a is the v8 generation’s headline result; v8b documents the failure mode that bounds it from above. Future work on the community head (Section 7) will not pursue further increases in supcon weight or temperature sharpening on this architecture, and will instead target either archetype-conditioned direct supervision (Section 5.8 hypothesis (a)) or a fundamentally different objective for orthogonalizing the trajectory manifold.

5.12 Cross-backbone signal probe (Qwen 2.5-7B / Qwen 3-8B / Mistral-7B-v0.3)

The single-backbone limitation flagged in §6.6 is partially addressable without re-training the adapter on a new backbone. The first question to ask is whether the discourse-community signal that v8a’s CDH amplifies is *specific to Qwen 2.5*, or whether it is latent in other 7B-class causal LMs and merely needs to be read out. We answer with the simplest possible probe: take raw mean-pooled hidden states from each backbone on `data/val_200.jsonl` (200 samples, 35 coarse Reddit communities), then compute leave-one-out 1-nearest-neighbor recall@1 on cosine similarity. No training, no labels at inference, no adapter.

Backbone	best raw-hidden recall@1	best layer	× chance
Qwen/Qwen2.5-7B	0.260	16	9.1×
Qwen/Qwen3-8B	0.220	4	7.7×
mistralai/Mistral-7B-v0.3	0.295	16	10.3×

For reference, chance on this 35-class probe is $1/35 = 0.0286$, and the v8a SRT-Adapter on Qwen/Qwen2.5-7B reaches recall@1 = 0.484 (16.7× chance) on the larger Reddit validation split (§5.9). Three observations follow.

The signal is not Qwen-specific. Mistral-7B-v0.3 raw hidden states actually carry slightly more discourse-community information than Qwen2.5-7B raw hid-

den states do ($10.3\times$ vs $9.1\times$ chance on the same probe). The community-discriminative subspace that the adapter’s CDH learns to amplify is present in both Qwen-family and Mistral-family backbones at comparable strength. Whatever the adapter is reading off the residual stream, the *raw substrate* it depends on is a property of contemporary 7B-class causal LMs, not of one specific tokenizer or training recipe.

The adapter’s contribution is a $\sim 1.8\times$ amplification of that substrate. v8a on Qwen 2.5-7B delivers $16.7\times$ chance, which is roughly $1.8\times$ the strongest raw-hidden-state baseline in the table (Mistral, $10.3\times$) and $1.8\times$ the same backbone’s own raw signal at the best layer ($9.1\times$). The adapter is doing real discriminative work on top of what the frozen backbone already encodes; the residual stream alone gets you most of the way to chance-level retrieval but does not reach the v8a number on its own.

The architecture should transfer. The CDH consumes a single hidden-state tensor at `community_layer_idx`; nothing in the adapter is Qwen-tokenizer-specific. Re-training the v8a recipe on Mistral-7B-v0.3 is a plausible future-work item; the raw-substrate result above is the strongest zero-cost evidence we currently have for the cross-backbone hypothesis. The full reproduction script is `scripts/cross_backbone_probe.py`, JSON results are at `artifacts/cross_backbone.json` and the two extra-backbone files (`cross_backbone_extra.json`, `cross_backbone_mistral.json`), and the rendered visualization is Fig. 8 of the public demo card.

5.13 Interiority study: post-release diagnostics on v8a

This subsection summarizes a sequence of post-release probes on the v8a checkpoint that dissect *how* the adapter encodes semiotic regime, with full numbers and reproduction commands in `docs/INTERIORITY_V1_FINDINGS.md`. The probe battery is 11 regimes (`metaphor`, `counterfactual`, `negation_modality`, `irony`, `literal`, `refusal_bait`, `lyric`, `deixis`, `self_reference`, `quoted_speech`, `code`) $\times$ 25 prompts. The community vector throughout is the 64-D CDH output (`use_prototypes=False`).

Layer-as-organ map. Per-regime z-score winners across nine adapter channels show channel specialization: `code` peaks on $\hat{r}$, regime entropy, and `div L21`; `metaphor` peaks on `inj L14`, `inj L21`, and the chain residual; `literal` peaks on `P(super)`. The four-module decomposition acquires functionally distinct channels under v8a’s continuous architecture.

Community 1-NN. Leave-one-out 1-NN top-1 accuracy on the 11-regime probe battery is 0.70, with bootstrap mean 0.77 and CI95 [0.72, 0.82] over 2000 iterations. This is $7.7\times$ chance on a 1/11 task. The top centroid pair `code` $\leftrightarrow$ `literal` (Euclidean $d = 1.07$) stays the most-separated pair in 85.5% of bootstrap iterations; the bottom pair `quoted_speech` $\leftrightarrow$ `self_reference` ($d = 0.39$) stays last in 83.0%. The point estimate of community separation ratio is 0.85, which we treat as a lower bound: the bootstrap CI95 [0.86, 0.98] is upward-biased by

within-cluster duplicate sampling.

BOS-register length scaling. A per-token trajectory probe across $T \in \{16, 49, 79, 129\}$ token windows (4 length bins $\times$ 5 prompts $\times$ 11 regimes) shows that the adapter’s write channels (inj L14, inj L21, div L7/L14/L21, chain residual) form a *fixed-size BOS register*: BOS-token amplitude on those channels is stable across an $8\times$ length sweep, staying within $\pm 2.5\%$ of the $T \approx 16$ baseline (peak-to-peak $\leq 4.8\%$). v8a learns to compress each prompt’s regime structure into a single fixed-amplitude write at the BOS slot; mid-prompt activity on inject channels is small.

BOS-sink ablation, per channel $\times$ regime. The above is consistent with two readings: either the adapter encodes its global regime summary at BOS and only at BOS, or BOS is a sink that contributes to but does not solely determine the regime ranking. We disambiguate by computing, for each channel c and regime g , the z-score of g ’s mean activation against the other ten regimes twice: once with BOS included in the pooling window ($z_{c,g}^{\text{with}}$), once with BOS excluded ($z_{c,g}^{\text{no}}$), and reporting $\Delta|z|_{c,g} = |z_{c,g}^{\text{with}}| - |z_{c,g}^{\text{no}}|$. On $\hat{r}$ and on regime entropy, $|\Delta|z| \leq 0.27$ for every regime: the regime-classifier scalars read essentially the same regime structure with or without the BOS slot, so the regime signature is *content-borne*. On inj L14 the deltas are large and signed: **metaphor** $\Delta = -0.64$ (regime signal is BOS-localized; removing BOS *strengthens* the relative z-score of the content-token mean), **literal** $\Delta = +0.98$ (regime signal is *anti*-localized to BOS; removing BOS dilutes the regime signature). The chain-residual channel is intermediate (max $|\Delta| = 0.58$ on **quoted_speech**). The honest synthesis is that the inject channels are functioning as register slots that hold a global per-prompt summary, while the regime classifier reads regime structure off content tokens directly. A simpler aggregate measure, regime-rank preservation under BOS exclusion, holds in 9/9 channels with mean Spearman $\rho = +0.74$, which is consistent with but coarser than the per-channel per-regime breakdown above.

Within-prompt trajectories. Plotting per-position curves for the regime-classifier scalars ($\hat{r}$, regime entropy) across 19 content position bins (BOS dropped) per regime shows that the temporal structure hidden by mean pooling is informative. $\hat{r}$ separates regimes by sustained amplitude: **code** sits at $\hat{r} \geq 0.85$ across the full window; **negation_modality**, **deixis**, and **refusal_bait** sit at $\hat{r} \leq 0.40$; the other seven regimes occupy a middle band 0.45–0.75. Regime entropy puts **code** in a sustained 0.18–0.32 band while every other regime stays at ≤ 0.05 across all 19 content positions. The classifier is consistently *uncertain* about the regime label whenever the prompt is **code** and consistently confident on every other regime, and that uncertainty is steady-state across positions, not a transient. Mean-pooled scalars collapse this to a single number per prompt and lose the structure.

Architectural reading of the interiority study. v8a learned to compress each prompt into a fixed-size *regime fingerprint* written into the BOS register on inject channels, while the regime classifier itself operates as a content-distributed

channel that reads regime evidence off non-BOS tokens. This is consistent with §6.3: the inject-back arm carries no measurable signal in v8a, so the adapter’s content-routing capacity is under-used. v9 (Section 7) targets this directly with a coverage-loss term that penalizes register concentration. The interiority probes do not change v8a’s headline numbers; they explain *where* in the architecture those numbers come from. Visualizations (per-channel z-score heatmaps, per-regime within-prompt trajectories, BOS-ablation $\Delta|z|$ matrix) appear as Figs. 10 and 11 of the public demo card.

6. Discussion

6.1 Semiotic Structure in Frozen Representations

The adapter architecture embodies a specific theoretical claim: that the semiotic structure of discourse, comprising community-conditioned interpretations, divergence patterns, and bifurcation dynamics, is already encoded in the hidden states of pretrained language models. The claim follows necessarily from the fact that these models were trained on text produced by communities with divergent interpretive norms. What the adapter adds is not new information but new *readout apparatus*: projections, attention mechanisms, and recurrence that disentangle the semiotic structure already present.

This is analogous to the relationship between a microscope and the structures it reveals. The adapter does not create bifurcation dynamics in text. It provides the lenses through which dynamics that were always present become visible and measurable.

6.2 Community Discovery Without Labels

A significant departure from the original SRT is the replacement of supervised community embeddings with unsupervised prototype-based clustering. The original architecture required explicit community IDs at both training and inference time, limiting deployment to domains with known community structure. The adapter’s community head learns to partition discourse space from backbone hidden states alone, discovering whatever grouping structure best serves the downstream semiotic losses.

This is more faithful to Peirce’s framework, in which communities of interpretation are not given *a priori* but emerge through shared interpretive practice. The prototypes are pulled apart by the semiotic losses: if assigning text to different communities helps the model predict divergence better, it will learn to separate them. Community structure is discovered, not imposed.

6.3 The Injection Pathway: Observation vs. Intervention

The RRM’s injection mechanism creates a feedback loop between semiotic observation and language generation. This is the architectural instantiation of metapragmatic awareness: the model’s observation of divergence changes the hidden states that produce subsequent text. The injection is deliberately small (scale factor $\alpha = 0.1$, zero-initialized projection, sigmoid gating), reflecting a conservative design philosophy: the adapter should primarily *observe* semiotic dynamics. Active intervention, that is, generation that *responds* to detected bifurcation, is an advanced capability that requires careful validation before scaling.

The CE loss provides a natural safety valve. Since gradients from CE flow through the injection pathway, the model is penalized if injections degrade language modeling quality. This creates an automatic pressure toward injections that are either helpful or neutral, never harmful.

The empirical situation through v8b is that this pressure has resolved on the *neutral* side of the helpful/neutral boundary: ablating the inject-back arm at evaluation produces no measurable downstream change on CE, on the regime classifier, or on the hallucination probes (Sections 2.5, 5.7, 5.9). The safety valve held, but the channel through which the meta-state would modulate generation has not learned to carry signal. The interiority study in §5.13 sharpens this picture: the inject channels *do* carry signal, but it is concentrated in a fixed-size BOS register (BOS amplitude is stable across an $8\times$ length sweep, peak-to-peak $\leq 4.8\%$), with the per-channel $\times$ per-regime BOS-ablation analysis showing that some regimes (e.g. **metaphor**) localize their inject signature at BOS and others (e.g. **literal**) are anti-localized to BOS. The arm has learned to *summarize* but not to *modulate*. v9’s coverage-loss design (Section 7, Appendix A version history) is directed at exactly this: penalize register concentration and credit distributed mid-prompt activity, so that the inject channel has gradient pressure to behave as a per-token modulator rather than a per-prompt summary slot.

This null is not new. Stage 3 Phase 1 ran 105 training rounds (R21 through R105 on TinyLlama-1.1B, documented in Lancaster, 2026a) with the explicit goal of activating the inject-back arm, including a remediation campaign covering loss-weight sweeps, BEN architecture overhauls, FiLM scale schedules, and gradient-isolation experiments. The arm did not activate on TinyLlama and has not activated through v8b on Qwen. Two readings are consistent with the data: (1) FiLM injection requires different scaffolding when the backbone is frozen, suggesting alternative injection mechanisms (cross-attention from RRM meta-state into selected backbone layers, low-rank residual-stream modulation conditioned on $\hat{r}$, learned gating bypassed by a meta-state classifier) are the appropriate v9-onward design target; (2) the mutual information carried by RRM meta-state about downstream gradients is fundamentally small when the backbone is frozen (per the information-theoretic bound in Section 2.6), placing a low ceiling on the intervention arm’s possible effectiveness. We do not have

the experimental record to distinguish these readings yet. v9 onward targets the first hypothesis directly. The current adapter should be characterized as a self-organizing observation channel over a frozen backbone, not yet as a closed circular-causal system.

6.4 Relation to the Pitchfork Model

BEN’s $\hat{r}$ estimate is the primary output of the entire system. It provides a per-token, continuous measure of semiotic stability that maps directly onto the control parameter of the pitchfork bifurcation:

- $\hat{r} < 0$: The sign is in the subcritical regime. Shared meaning is stable. Perturbations decay.
- $\hat{r} \approx 0$: The sign is near-critical. Small changes in context or community could tip it.
- $\hat{r} > 0$: The sign has bifurcated. Meaning has split into community-specific attractors.

This is not a classifier applied after the fact. $\hat{r}$ is estimated from the accumulated meta-state of the RRM, which tracks how divergence has evolved through the backbone’s processing hierarchy. It is a real-time structural estimate, not a post-hoc label.

6.5 What $\hat{r}$ actually measures

The context-conditional probe in §5.6 returned a null result at the target token and a *negative* result over the full passage. This is informative. Earlier drafts of this paper described $\hat{r}$ as a contestedness detector, that is, as a per-token estimate of how much a sign is being fought over in its current discourse register. The data does not support that interpretation.

What $\hat{r}$ actually appears to measure is information density combined with the specific reflexivity components encoded in r_{true} (political-lean magnitude, annotator divergence, connection density). Fact-dense prose (dates, numbers, named entities, citations) drives $\hat{r}$ up because those are the positions where r_{true} is highest in the training distribution. Rhetorical or formulaic charged language is, on average, *less* lexically diverse than dense factual writing, so its mean $\hat{r}$ is lower.

The contestedness signal is in the *community head*, not in $\hat{r}$. The counterfactual-decoding result (§5.3) and the per-topic community shifts in §5.6 both demonstrate this. The v6 divergence-SupCon loss was intended to sharpen the metapragmatic channel further; in practice it improved community recall@1 (§5.7) but did not produce the expected community-conditional separation of MAH divergence trajectories on contested topics, and v8a’s prototype-bottleneck removal (§5.9) turned out to be the more consequential change for both channels.

This revision of the architectural story is cleaner, not weaker: the model has two distinct outputs that measure two distinct things, and the data tells us which

is which.

6.6 Limitations

1. **No modulation at inference.** The current architecture estimates $\hat{r}$ but does not use it to modulate generation. Future work will explore λ -controlled modes where detected bifurcation triggers bridge-generation strategies.
2. **Simplified regime model.** The binary subcritical/supercritical classification omits the near-critical regime, which is arguably the most important for practical applications (early warning of emerging bifurcation). The three-class model from the original SRT will be restored once binary classification is validated.
3. **Reddit-only data.** Training on Reddit discourse may not generalize to other domains (news media, academic text, legal documents). The TruthfulQA hallucination probe (§5.4) is a partial cross-domain transfer signal but mean $\hat{r}$ AUROC of 0.573 is well below the 0.7 production threshold.
4. **No human evaluation.** All supervision comes from computed r_{true} labels. Ecological validity, that is, whether $\hat{r}$ tracks what human annotators perceive as meaning contestation, has not been tested. The §5.6 negative result is the strongest current evidence that r_{true} as currently constructed is not a clean proxy for contestedness.
5. **Single backbone (partially answered).** Adapter-trained results are reported only for Qwen 2.5-7B. The cross-backbone raw-hidden probe in §5.12 shows the discourse-community signal that v8a amplifies is present at comparable strength in Qwen/Qwen3-8B ($7.7\times$ chance on raw hidden states) and mistralai/Mistral-7B-v0.3 ($10.3\times$ chance, slightly stronger than Qwen 2.5-7B’s own $9.1\times$). This is consistent with the architecture being backbone-agnostic and with the adapter’s $\sim 1.8\times$ amplification factor transferring, but the actual adapter-on-Mistral re-train has not yet been executed. The full backbone-agnostic claim still requires that re-training, which is a v9-onward work item.

6.7 Connections to Emergent Perspective Diversity in Reasoning Models

A separate line of recent work has documented that language models trained with reasoning-style reinforcement spontaneously develop heterogeneous internal features that mechanistic interpretability methods can read out as something like distinct personalities, areas of expertise, or stances, and that these features enter into structured conflict and reconciliation during chain-of-thought generation (Evans, in preparation; cf. Foster, Rzhetsky, & Evans, 2015 for the cross-corpus precedent). The relationship between that finding and the present work is structural rather than methodological. Evans’s program reveals what emerges

spontaneously under reasoning supervision in an architecturally-undifferentiated backbone; the SRT-Adapter provides explicit architectural channels (the four Peircean subspaces of the semiotic embedding layer in the full SRT, and the community / divergence / bifurcation subspaces of the present adapter) into which similar emergent structure can organize. The two approaches are complementary in a specific sense: probing the unstructured backbone reveals the existence of perspective-like features without a vocabulary for what they are; the adapter’s explicit decomposition supplies a vocabulary, grounded in Peircean semiotics and linguistic anthropology, but only weakly constrains what fills the slots.

This suggests a research question that neither approach can resolve alone. Reasoning models that develop internal perspective diversity may be performing genuine semiotic work, namely navigating meaning divergence across implicit interpretive communities, or they may be performing something closer to cognitive brainstorming or rhetorical variation within a single interpretive framework. The SRT-Adapter’s metapragmatic attention head is designed to detect the former and would, in principle, register low divergence across the latter. Cross-applying mechanistic interpretability tools to adapter-equipped reasoning models, and cross-applying the adapter’s divergence and bifurcation readouts to backbones probed for emergent features, is the form of collaboration this paper points toward as the next step.

The von Foerster framing (Section 2.5) makes the same point in cybernetic terms: spontaneous perspective diversity in reasoning models is self-organization without an explicit semiotic substrate; the adapter provides the substrate without yet exhibiting the fully closed circular-causal loop. The two findings, taken together, suggest that the substrate and the emergent dynamics may be separable engineering targets, and that the productive question is what becomes possible when both are present.

7. Conclusion

The SRT-Adapter demonstrates that semiotic awareness can be added to any frozen language model as a lightweight, modular capability. By tapping hidden states rather than rebuilding the backbone, the architecture preserves pretrained language modeling quality (CE = 2.63 vs. unadapted 2.71 on the same val data) while introducing structured outputs that make the semiotic dynamics of text visible and measurable.

The v5 generation establishes the basic capability set on five independent probes:

1. **CE preservation** (§5.1): the injection pathway is mildly helpful, not harmful.
2. **Community geometry** (§5.2): recall@1 of 0.36 on a 35-class unsupervised retrieval task (12.6× random), made possible by SupCon on the

encoder’s pre-mixing output.

3. **Counterfactual community decoding** (§5.3): forcing the community vector at decode time produces zero disagreement on factual prompts and near-total disagreement on contested topics, demonstrating the community vector behaves as a discourse prior.
4. **Hallucination signal** (§5.4): all four SRT-native channels lean in the predicted direction on TruthfulQA without truthfulness supervision; mean $\hat{r}$ AUROC = 0.573.
5. **Regime calibration** (§5.5): ECE = 9×10^{-4} and AUROC = 0.99 on 351K tokens, unblocking downstream probabilistic use.

The negative result on context-conditional $\hat{r}$ (§5.6) sharpened the architectural story: $\hat{r}$ measures information density and the components of r_{true} , while the community head measures discourse register. Conflating these in earlier drafts was a theoretical error the data corrected.

The v6–v8 generations then reframe the design. v6 and v7 (§5.7) extend SupCon from the community channel to the metapragmatic divergence channel and add ListNet ranking on $\hat{r}$, giving incremental gains on community recall@1 (0.360 $\rightarrow$ 0.413) and calibration (ECE $\rightarrow$ 0.0006). The cross-corpus convergence probe (§5.8) shows the prototype tensors barely move during training and that 33 externally-curated archetypes collapse onto roughly four functional macro-clusters, reproducing at single-backbone scale the macro-attractor pattern Foster, Rzhetsky, & Evans (2015) document for scientific subdisciplines. v8a (§5.9) is the headline architectural result: removing the 32-prototype mixing layer entirely leaves CE unchanged ($\Delta = +0.0001$ nats) while raising Reddit recall@1 from 0.413 to 0.484, raising archetype recall@1 to $7.6\times$ chance, nearly doubling the within/between cosine ratio, and expanding trajectory anisotropy by $\sim 325\times$. v8b (§5.10) falsifies the “sharper-supcon” hypothesis on the continuous-encoder architecture, bounding the v8 design from above. The encoder, not the prototype basis, was doing the discriminative work; the prototype layer was discarding it through a saturated soft-argmax.

What the v5–v8 arc has demonstrated is a *self-organizing observation channel over a frozen backbone* (Sections 2.5–2.6, 5.1–5.10). What it has not yet demonstrated is a *closed circular-causal loop* in which the meta-state’s observation modifies generation in a measurable way: ablating the inject-back arm produces no downstream change on CE, calibration, or hallucination probes through v8b (Sections 2.5, 6.3). Reading this through the information-bound framing of Section 2.6, the simplest hypothesis is that the mutual information the meta-state currently carries about the downstream loss is small, and that pushing Q_0 above the bifurcation threshold of the inject-back arm requires either a larger meta-state, a different injection geometry, or direct supervision on the closed-loop behavior. v9 onward is the design target for that work, together with archetype-conditioned direct supervision (§5.8 hypothesis (a)) for resolving sub-macro-cluster archetype structure. Combined-feature hallucination probes, cross-domain transfer evaluation, and human ecological-validity studies remain

the principal open empirical questions.

References

- Agha, A. (2003). The social life of cultural value. *Language & Communication*, 23(3–4), 231–273.
- Anderson, M. (2014). Mathematical modeling of catastrophic change in cultural systems. In M. Anderson (Ed.), *Cultural shaping of violence: Victimization, escalation, response* (selected chapters). Purdue University Press.
- Bail, C. A., et al. (2018). Exposure to opposing views on social media can increase political polarization. *Proceedings of the National Academy of Sciences*, 115(37), 9216–9221.
- Bennett, C. H. (1982). The thermodynamics of computation: A review. *International Journal of Theoretical Physics*, 21(12), 905–940.
- Deely, J. (2014). The suprasubjective in semiotic relations. *The American Journal of Semiotics*, 30(3–4), 165–182.
- Durst-Andersen, P. (2011). *Linguistic supertypes: A cognitive-semiotic theory of human communication*. De Gruyter Mouton.
- Evans, J. A. (2010). Industry induces academic science to know less about more. *American Journal of Sociology*, 116(2), 389–452.
- Foster, J. G., Rzhetsky, A., & Evans, J. A. (2015). Tradition and innovation in scientists’ research strategies. *American Sociological Review*, 80(5), 875–908.
- Irvine, J. T., & Gal, S. (2000). Language ideology and linguistic differentiation. In P. V. Kroskrity (Ed.), *Regimes of language* (pp. 35–83). SAR Press.
- Kockelman, P. (2017). *The art of interpretation in the age of computation*. Oxford University Press.
- Kockelman, P. (2024). *Last words: A theory of everything that matters*. University of Chicago Press.
- Kockelman, P. (2025). *Semiotic agency in digital environments*. Manuscript.
- Lancaster, J. B. (2025). The treachery of signs: Semiotic mediation, pitch-fork bifurcation, and political polarization in algorithmically curated societies. SSRN. <https://papers.ssrn.com/abstract=5987495>
- Lancaster, J. B. (2026a). Semiotic-reflexive language model training: Bridging interpretive bifurcations through metapragmatic chain architectures and embodied grounding. SSRN. <https://papers.ssrn.com/abstract=6349978>
- Lancaster, J. B. (2026b). Prenatal origins of cross-modal iconic correspondence: A semiotic analysis.

- Lancaster, J. B. (2026c). Reddit Discourse Corpus: A multi-community dataset for semiotic analysis.
- Landauer, R. (1961). Irreversibility and heat generation in the computing process. *IBM Journal of Research and Development*, 5(3), 183–191.
- Latour, B. (1996). On interobjectivity. *Mind, Culture, and Activity*, 3(4), 228–245.
- Leighton, M. P. (2026). Will a large complex system be a Maxwell demon? *arXiv preprint arXiv:2603.03248*.
- Mangalam, M. (2025). Against the Bayesian brain. *Behavioral and Brain Sciences* (forthcoming).
- Maturana, H. R., & Varela, F. J. (1980). *Autopoiesis and cognition: The realization of the living*. D. Reidel.
- Parrondo, J. M. R., Horowitz, J. M., & Sagawa, T. (2015). Thermodynamics of information. *Nature Physics*, 11(2), 131–139.
- Peirce, C. S. (1931–1958). *Collected papers of Charles Sanders Peirce* (Vols. 1–8). C. Hartshorne, P. Weiss, & A. Burks (Eds.). Harvard University Press.
- Radford, A., et al. (2021). Learning transferable visual models from natural language supervision. In *ICML 2021*.
- Ramachandran, V. S., & Hubbard, E. M. (2001). Synaesthesia: a window into perception, thought and language. *Journal of Consciousness Studies*, 8(12), 3–34.
- Silverstein, M. (1993). Metapragmatic discourse and metapragmatic function. In J. A. Lucy (Ed.), *Reflexive language* (pp. 33–58). Cambridge University Press.
- Silverstein, M. (2003). Indexical order and the dialectics of sociolinguistic life. *Language & Communication*, 23(3–4), 193–229.
- VanSaders, B., Fruchart, M., & Vitelli, V. (2026). Measurement-induced phase transitions in informational active matter. *PNAS Nexus*, pgag077. <https://doi.org/10.1093/pnasnexus/pgag077>
- Versace, E., et al. (2023). Cross-modal correspondences between auditory and visual features in domestic chicks. *Animal Cognition*, 26, 1021–1030.
- von Foerster, H. (1981). *Observing systems*. Intersystems Publications.
- von Foerster, H. (2003). *Understanding understanding: Essays on cybernetics and cognition*. Springer.
- Wildgen, W. (1982). *Catastrophe-theoretic semantics: An elaboration and application of René Thom’s theory*. John Benjamins.

Appendix A: Configuration Defaults (v8a)

```

SRTConfig(
    backbone_id      = "Qwen/Qwen2.5-7B",
    backbone_dtype   = "bfloat16",
    mah = MAHConfig(d_sub=512, d_divergence=256, num_heads=4, dropout=0.1),
    rrm = RRMConfig(d_meta=512, inject_scale=1.0),    # FiLM since v4
    ben = BENConfig(d_hidden=256),                  # tanh removed in v4
    community = CommunityConfig(
        num_prototypes=32, d_community=64, temperature=1.0,
        use_prototypes=False,                        # v8a: encoder output IS the community
    ),
    loss = LossConfig(
        ce_weight=1.0, chain_weight=0.5, bif_weight=1.0,
        regime_weight=5.0, div_alive_weight=0.1,
        inject_reg_weight=0.0, inject_target_norm=1.0,    # v4: dropped (FiLM init handles
        community_entropy_weight=0.01,
        # v5
        community_supcon_weight=2.0,
        community_supcon_temperature=0.1,
        # v6
        divergence_supcon_weight=0.3,                    # v7: dropped from 1.0 to recover
        divergence_supcon_temperature=0.1,
        listnet_weight=0.5,
        listnet_temperature=1.0,
        chain_residual_aux_weight=0.05,
        chain_residual_aux_target=0.5,
    ),
)

```

Version history

Version	Headline change	Result
v1	Initial architecture, tanh on $\hat{r}$, L2 on injections	$\hat{r}$ saturated at ± 1 , injections at norm ≈ 7
v2	Diagnostic instrumentation added	Confirmed both pathologies; CE healthy
v3	Target-norm injection penalty $(\ \text{inj}\ - 1)^2$	Injections recovered; community prototypes still collapsed (recall@1 = 0.05)

Version	Headline change	Result
v4	tanh removed from BEN; RRM linear-gated $\rightarrow$ FiLM; first SupCon attempt on vector (failed)	$\hat{r}$ tail recovered; SupCon flatlined at $\log(B - 1)$
v5	SupCon switched to encoded (pre-mixing); weight $0.5 \rightarrow 2.0$; warm-restart of community head	recall@1 = 0.36, ECE = 9×10^{-4} , counterfactual-decode contested/factual split
v6	Add divergence-SupCon ($\lambda = 1.0$), ListNet on $\hat{r}$, chain-residual auxiliary floor	Reddit recall@1 $0.36 \rightarrow 0.41$, ECE $\rightarrow 0.0006$; §5.3 counterfactual decode regressed
v7	Reduce divergence-SupCon to $\lambda = 0.3$	Reddit recall@1 $\rightarrow 0.413$, hallu AUROC $\rightarrow 0.5785$, decode signal recovered
v8a	Drop the 32-prototype mixing layer (<code>use_prototypes=False</code>), encoder output = community vector	CE $\Delta = +0.0001$; Reddit recall@1 $\rightarrow 0.484$; archetype recall@1 $7.6\times$ chance; within/between cosine ratio $\rightarrow 2.016$; trajectory anisotropy $\times 325$
v8b	Sharpen community-SupCon ($\lambda: 2.0 \rightarrow 4.0$, $\tau: 0.10 \rightarrow 0.05$) on v8a base	Partial regression on every encoder-geometry metric except anisotropy; falsifies “sharper is better” on this architecture

Version	Headline change	Result
v8a interiority (post-release)	Eight diagnostic probes on the v8a checkpoint (§5.13): layer-as-organ map, community 1-NN, BOS-register length scaling, BOS-sink ablation per channel $\times$ regime, within-prompt trajectories, regime calibration replication, cross-backbone raw-hidden probe (§5.12)	Confirms v8a inject channels function as a fixed-size BOS register (peak-to-peak $\leq 4.8\%$ across $8\times$ length sweep); regime classifier reads content-borne signal; calibration ECE replicates at 9.1×10^{-4} on 351K tokens; community 1-NN at $7.7\times$ chance; cross-backbone signal latent in Qwen 3-8B and Mistral-7B at comparable strength
v9 (in progress)	Closed-loop training target for the inject-back arm (Sections 2.5, 6.3); coverage loss penalizing register concentration (§5.13, §6.3); archetype-conditioned direct supervision (§5.8 hypothesis (a)); adapter re-train on Mistral-7B-v0.3 (§6.6)	TBD

Appendix B: Layer Index Auto-Computation

Given backbone depth L : - MAH hook layers: $\lfloor L/4 \rfloor, \lfloor L/2 \rfloor, \lfloor 3L/4 \rfloor$ - RRM injection layers: MAH layers 2 and 3 (skip first to let meta-state accumulate) - Community discovery layer: $\max(1, \lfloor L/7 \rfloor)$

For Qwen 2.5-7B ($L = 28$): MAH @ [7, 14, 21], inject @ [14, 21], community @ 4.